



Advanced Matchlib Topics 2021

Stuart Swan | Platform Architect
Siemens EDA, a part of Siemens Digital Industries Software

Catapult SystemC / Matchlib Resources

This presentation assumes you have reviewed the slides and examples below:

Matchlib training slides

- `$MGC_HOME/shared/examples/matchlib/toolkit/doc/matchlib_customer_training.ppt`

Matchlib example kit in Catapult install

- `$MGC_HOME/shared/examples/matchlib/toolkit`

Brief Review of Matchlib Interfaces

- Message Passing Interfaces : Push/Pop, PushNB/PopNB
 - \$MGC_HOME/shared/examples/matchlib/toolkit/examples/05_push_pop
- Data + Valid Protocol :
 - ...toolkit/examples/15_data_valid_protocol
- Event Pulse :
 - ...toolkit/examples/94_event_pulse
- Process Sync:
 - ...toolkit/examples/12_ping_pong_mem
- External Dual Port RAM:
 - ...toolkit/examples/12_ping_pong_mem
- ARM AXI4:
 - ...toolkit/examples/08_dma
- ARM APB:
 - ...toolkit/examples/52_apb

Brief Review of Matchlib Interfaces (continued)

- Native SystemC signals :
 - ...toolkit/examples/04_ready_valid
- Toggle Protocol :
 - ...toolkit/examples/13_toggle_protocol
- FIFO instantiation between modules :
 - ...toolkit/examples/10_fifo_channel_hier
- FIFO instantiation between processes within a module :
 - ...toolkit/examples/11_fifo_channel_flat

Brief Review of Primary Matchlib Features

- Throughput accurate modeling in Pre-HLS simulation
 - ...toolkit/examples/06_thruput_accurate
- Waveform generation in pre-HLS simulation
 - ...toolkit/examples/06_thruput_accurate
- Channel log generation in pre-HLS simulation with auto-comparison utilities
 - ...toolkit/doc/matchlib_soc_debug_tutorial.pdf
- Supports multiple clocks, sync/async resets with automatic consistency checking
 - ...toolkit/examples/95_multi_clk
- Mixed language simulation and debug support in Questa, VCS, and Xcelium
 - ...toolkit/examples/45_vlog_tb_dma_dut

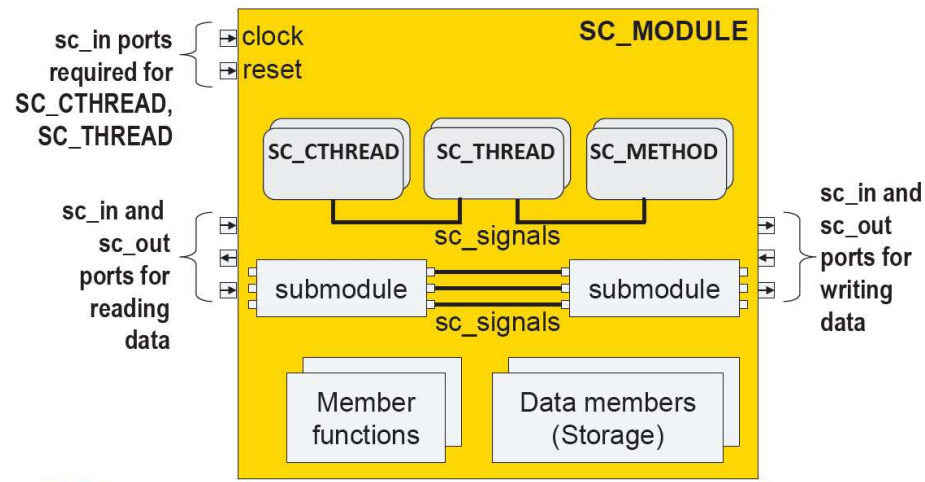
Brief Review of Primary Matchlib Features (continued)

- ~30X RTL simulation speed in ACCURATE_SIM mode
- ~300X RTL simulation speed in FAST_SIM mode
- Random stall injection in pre-HLS model (Matchlib DUT and TB.)
 - ...toolkit/doc/matchlib_soc_debug_tutorial.pdf
- Random stall injection in post-HLS model (Catapult RTL via STALL_FLAG_SV)
- Thruput and latency backannotation
 - ...toolkit/examples/71_annotate_simple

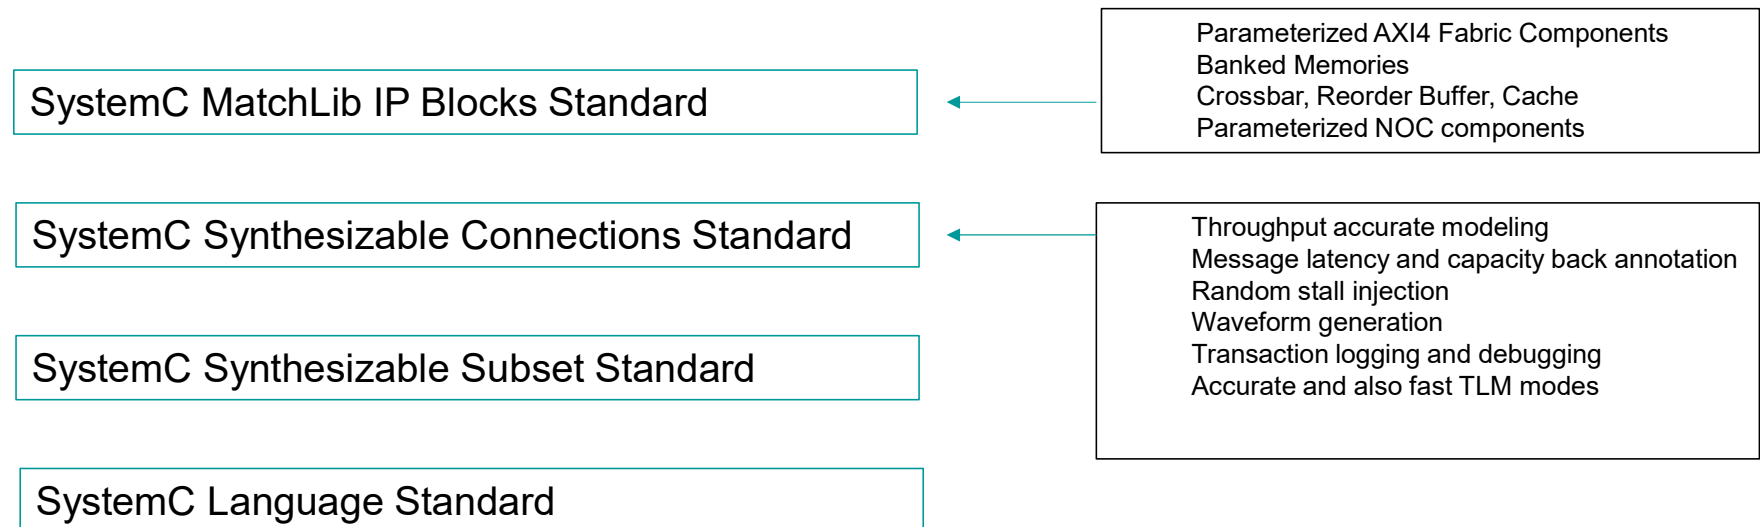
MatchLib Relationship to SystemC Standards

Accellera SystemC Synthesizable Subset Standard focuses on what's in diagram below

- Modules, Ports, Processes, clocks, resets, signal IO, datatypes



Proposed SystemC HLS Standards Layers



Matchlib Key Coding Guidelines: When do you need to use Non-blocking IO?

- Arbitration requires Peek or PopNB to all arbitrated inputs.
- Time-based splitting and merging of transaction streams requires PushNB and PopNB (respectively)
- Even when using NB IO, it is best to have **at least one blocking IO call** in each iteration of a loop
 - This ensures Catapult understands a clear relationship between II of loop and IO behavior.
- There are cases where a process will need ALL non-blocking IO (all PopNB and PushNB), but it should be pretty rare.
- In this case the process should be kept as small / simple as possible, ideally communicating with other processes that follow the guidelines above.
- With all non-blocking IO, you will likely be modeling at very close to RTL level, and most likely HLS will just be translating SC RTL into Verilog RTL.

Matchlib: Always prefer to use blocking IO over nonblocking IO.

- Your models will be simpler and more likely to have a good process structure.
- 100% blocking IO is called KPN (Kahn Process Networks)
 - KPN is deterministic
 - easier to verify.
- Non-blocking IO is sometimes needed, but introduces timing dependent behavior, and can make verification more difficult in some cases if the timing dependent behavior is externally visible.
- Summary: always ask yourself if use of non-blocking IO is really justified.

Key Guidelines for Coding for Good QOR in SC

- Simulate and debug your pre-HLS model before you synthesize (!)
 - Verify performance and functionality in your pre-HLS model
- In SystemC HLS, you should always have:
 - At least a **rough idea** of what your HW implementation will be (e.g. pipeline characteristics)
 - An **exact idea** of what your module pin level interfaces are
- **Refine your architecture in your pre-HLS model**
 - **Usually this is by far the most effective way to improve QOR**
- If functionality can easily be split into smaller processes, it is usually better to do so
 - Control FSMs generated by Catapult will be smaller
 - Smaller processes may be able to run in parallel

Can I use “feature X” with Catapult SystemC Flow?

- AC Datatypes – Yes
- AC Math – Yes
- AC DSP – Functions only, no hierarchy or instantiated AC channels
- AC channel – No
- “Interface synthesis” – Generally, No
- “Interface synthesis for C arrays in single process” - Yes
- Ccores – Yes
- Blackboxes – Yes
- Catapult Directives – Generally, Yes
- Coupled_IO - Yes

Matchlib Accuracy

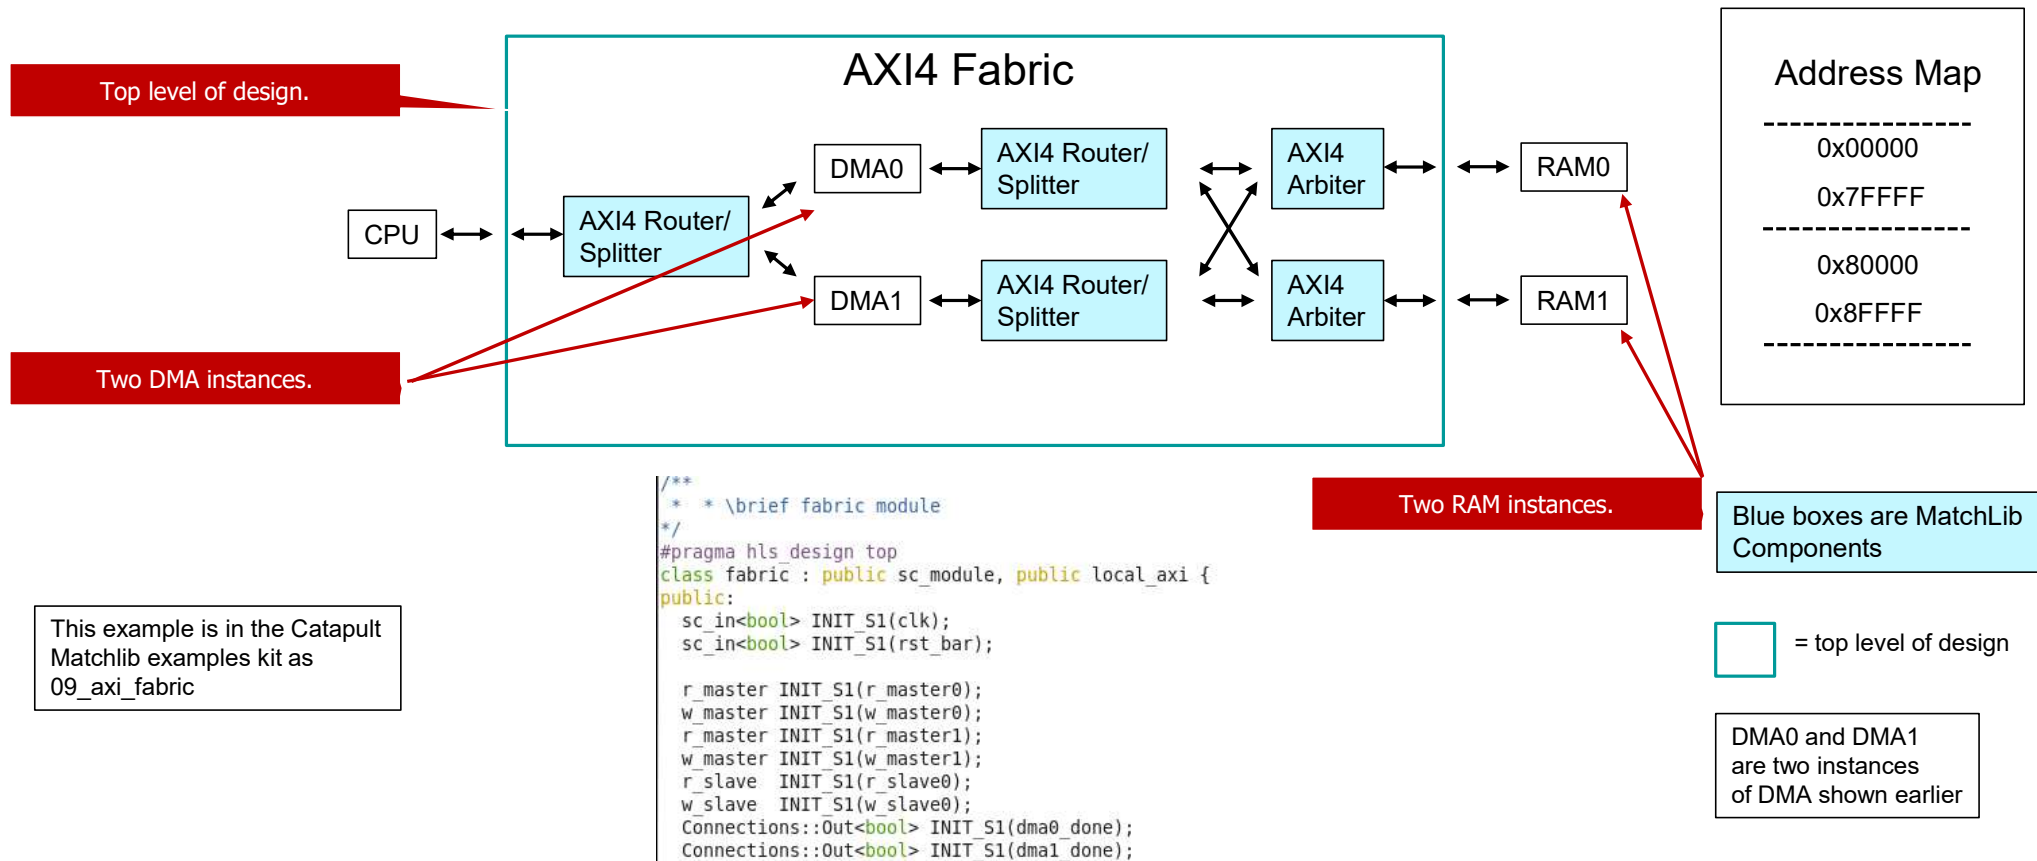
- Goal of Matchlib is to be “throughput accurate” in pre-HLS models, even as you scale up size of systems.
- Goal is NOT to be precisely cycle accurate wrt RTL
- Even with some differences in accuracy wrt RTL, there is big value in being able to model time/performance in the pre-HLS models.
- It is useful to understand where sources of inaccuracies may arise and how to address them if desired.

Small Example

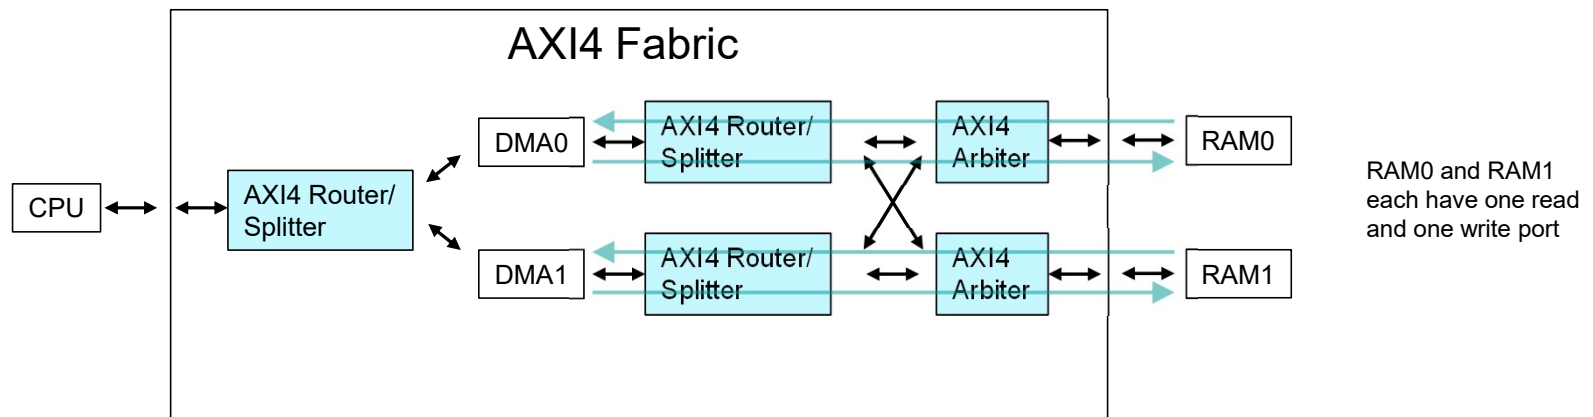
- Matchlib assumes all processes are pipelined with an $II=1$.
- Matchlib assumes `stall_mode=flush`.
- Matchlib assumes `coupled_io=false`.
- Matchlib assumes pipeline latency = 1
- Pre-HLS sim of model on right will Pop and Push one transaction on every clock.
 - Output will appear one clock after input.
- Post-HLS RTL will have same behavior except latency will increase.
 - `ac_sqrt` likely to require several clock cycles latency in RTL.

```
12 Connections::Out<uint32> CCS_INIT_S1(out1);
13 Connections::In <uint32> CCS_INIT_S1(in1);
14
15 SC_CTOR(dut) {
16     SC_THREAD(main);
17     sensitive << clk.pos();
18     async_reset_signal_is(rst_bar, false);
19 }
20
21 private:
22
23 void main() {
24     out1.Reset();
25     in1.Reset();
26     wait();
27 #pragma hls_pipeline_init_interval 1
28 #pragma pipeline_stall_mode flush
29     while (1) {
30         uint32_t i = in1.Pop();
31         uint32_t o = ac_sqrt(i);
32         out1.Push(o);
33     }
34 }
35 };
36
```

Larger Example: AXI4 Bus Fabric using MatchLib



AXI4 Bus Fabric using MatchLib – Test #0



Test #0: Concurrently,
DMA0 reads/writes 320 beats to RAM0
DMA1 reads/writes 320 beats to RAM1

(a beat is the data chunk transferred on the bus in 1 clock cycle)

AXI4 Bus Fabric Test #0 simulation logs

BEFORE HLS (SystemC simulation)


```
0 s top Stimulus started
6 ns top Running FABRIC_TEST # : 0
44 ns top.ram0 ram read  addr: 000000000 len: 0ff
44 ns top.ram0 ram write addr: 000002000 len: 0ff
49 ns top.ram1 ram write addr: 000002000 len: 0ff
49 ns top.ram1 ram read  addr: 000000000 len: 0ff
304 ns top.ram0 ram read  addr: 000000800 len: 03f
309 ns top.ram1 ram read  addr: 000000800 len: 03f
311 ns top.ram0 ram write addr: 000002800 len: 03f
316 ns top.ram1 ram write addr: 000002800 len: 03f
385 ns top dma_done detected. 1 1
385 ns top start_time: 46 ns end_time: 385 ns
385 ns top axi beats (dec): 320
385 ns top elapsed time: 339 ns
385 ns top beat rate: 1059 ps
385 ns top clock period: 1 ns
425 ns top finished checking memory contents
```

AFTER HLS (Verilog RTL simulation)

```
# 0 s top Stimulus started
# 6 ns top Running FABRIC_TEST # : 0
# 55 ns top/ram0 ram write addr: 000002000 len: 0ff
# 60 ns top/ram1 ram write addr: 000002000 len: 0ff
# 68 ns top/ram0 ram read  addr: 000000000 len: 0ff
# 70 ns top/ram1 ram read  addr: 000000000 len: 0ff
# 340 ns top/ram0 ram write addr: 000002800 len: 03f
# 342 ns top/ram1 ram write addr: 000002800 len: 03f
# 343 ns top/ram0 ram read  addr: 000000800 len: 03f
# 345 ns top/ram1 ram read  addr: 000000800 len: 03f
# 414 ns top dma_done detected. 1 1
# 414 ns top start_time: 55 ns end_time: 414 ns
# 414 ns top axi beats (dec): 320
# 414 ns top elapsed time: 359 ns
# 414 ns top beat rate: 1122 ps
# 414 ns top clock period: 1 ns
# 454 ns top finished checking memory contents
```

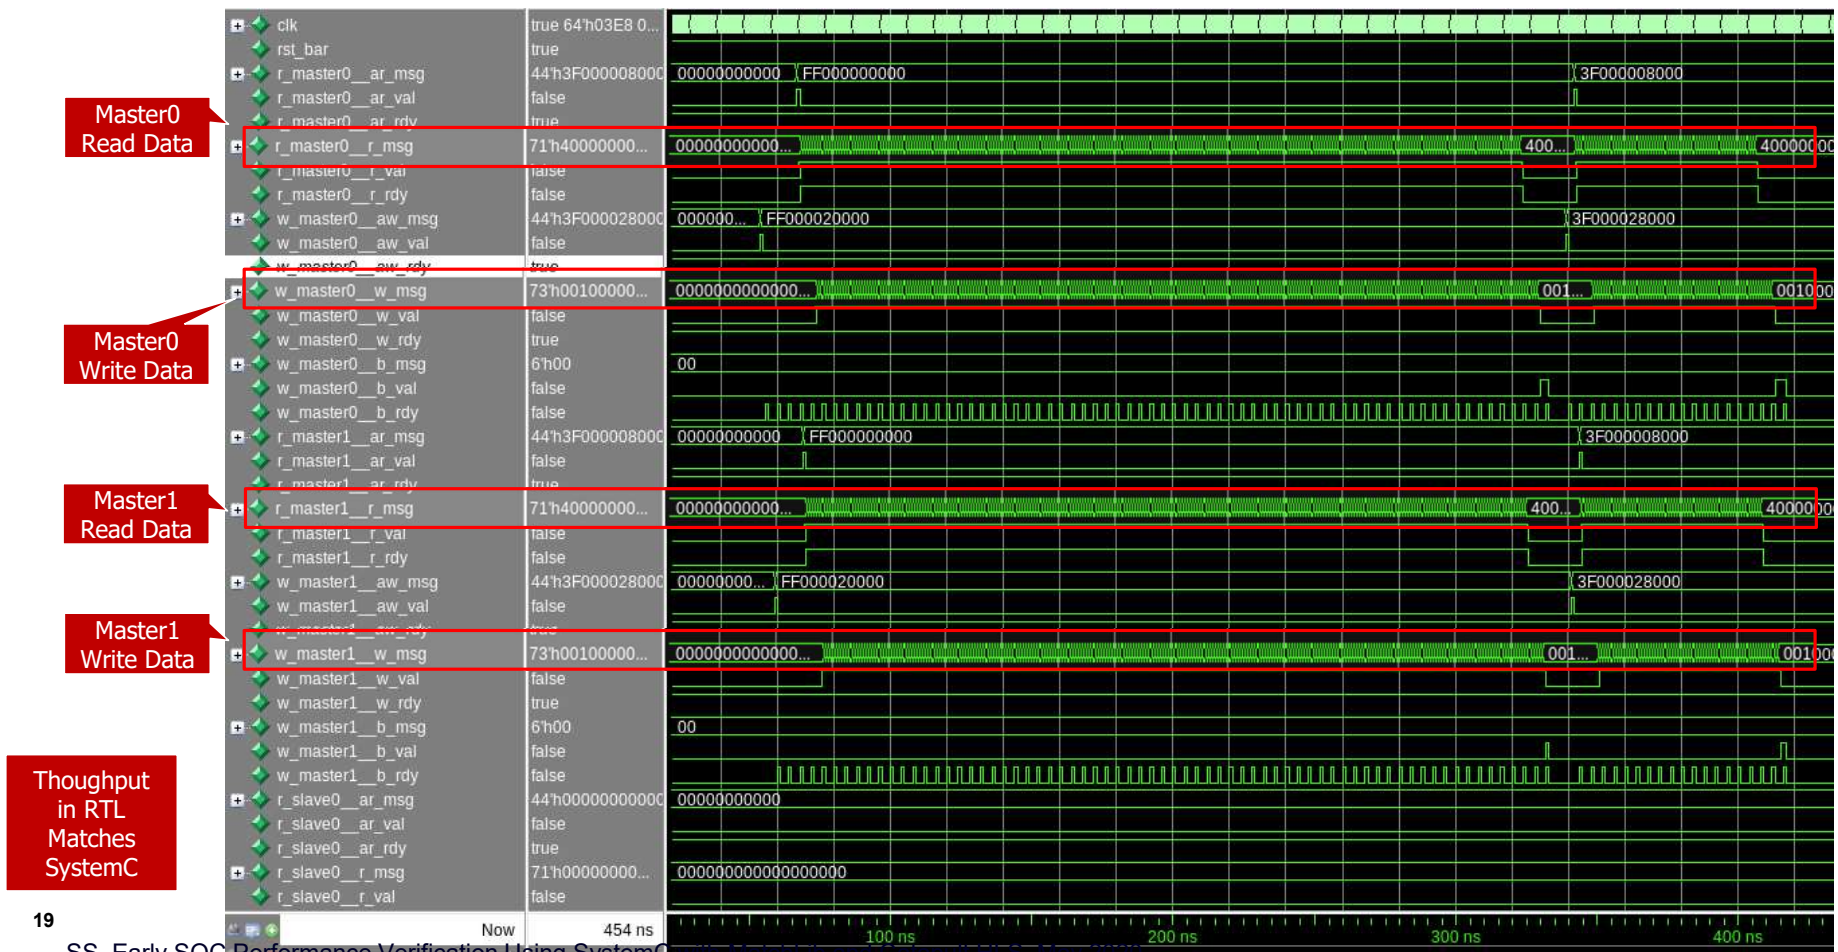
Before and after HLS we get nearly one beat per clock cycle

AXI4 Fabric Waveforms Before Catapult HLS–Test #0 (SystemC)

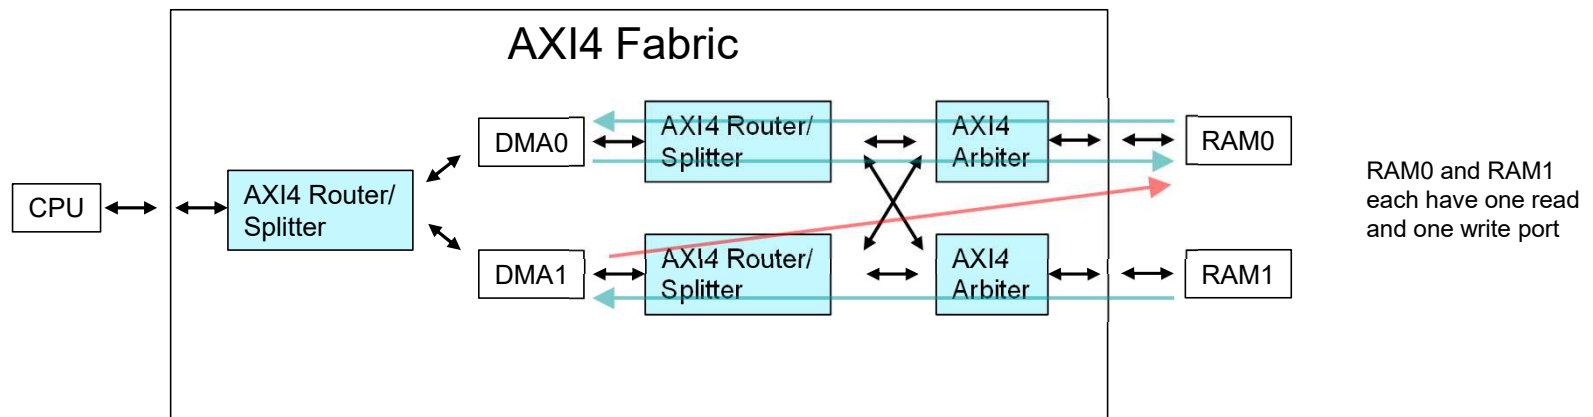


18

AXI4 Fabric Waveforms After Catapult HLS – Test #0 (Verilog)



AXI4 Bus Fabric using MatchLib – Test #1



Test #1: Concurrently,
DMA0 reads/writes 320 beats to RAM0
DMA1 reads 320 beats from RAM1 and writes to RAM0
Note contention on RAM0 writes

AXI4 Bus Fabric Test #1 simulation logs

BEFORE HLS (SystemC simulation)


```
0 s top Stimulus started
6 ns top Running FABRIC_TEST # : 1
44 ns top.ram0 ram read  addr: 000000000 len: 0ff
44 ns top.ram0 ram write addr: 000002000 len: 0ff
49 ns top.ram1 ram read  addr: 000000000 len: 0ff
304 ns top.ram0 ram read  addr: 000000800 len: 03f
308 ns top.ram0 ram write addr: 000006000 len: 0ff
560 ns top.ram1 ram read  addr: 000000800 len: 03f
566 ns top.ram0 ram write addr: 000002800 len: 03f
632 ns top.ram0 ram write addr: 000006800 len: 03f
701 ns top dma_done detected. 1 1
701 ns top start_time: 46 ns end_time: 701 ns
701 ns top axi beats (dec): 320
701 ns top elapsed time: 655 ns
701 ns top beat rate: 2047 ps
701 ns top clock period: 1 ns
741 ns top finished checking memory contents
```

AFTER HLS (Verilog RTL simulation)

```
# 0 s top Stimulus started
# 6 ns top Running FABRIC_TEST # : 1
# 55 ns top/ram0 ram write addr: 000002000 len: 0ff
# 68 ns top/ram0 ram read  addr: 000000000 len: 0ff
# 70 ns top/ram1 ram read  addr: 000000000 len: 0ff
# 335 ns top/ram0 ram write addr: 000006000 len: 0ff
# 343 ns top/ram0 ram read  addr: 000000800 len: 03f
# 598 ns top/ram1 ram read  addr: 000000800 len: 03f
# 598 ns top/ram0 ram write addr: 000002800 len: 03f
# 670 ns top/ram0 ram write addr: 000006800 len: 03f
# 736 ns top dma_done detected. 1 1
# 736 ns top start_time: 55 ns end_time: 736 ns
# 736 ns top axi beats (dec): 320
# 736 ns top elapsed time: 681 ns
# 736 ns top beat rate: 2128 ps
# 736 ns top clock period: 1 ns
# 776 ns top finished checking memory contents
```

Two concurrent writes to RAM0 cause beat rate to be above two clock cycles.

AXI4 Fabric Waveforms Before Catapult HLS –Test#1 (SystemC)



AXI4 Fabric Waveforms After Catapult HLS – Test #1 (Verilog)



23

SS, Early SOC Performance Verification Using SystemC with MatchLib and Catapult HLS, May 2020

Key Observations

- “Throughput accuracy” of Matchlib pre-HLS models is generally retained even as you scale up size of systems.
- Ability to view, analyze, and debug time-based behaviors and waveforms in pre-HLS models is very valuable to:
 - HW architects
 - HLS engineers
 - DV engineers
 - RTL designers on project who will “never touch” HLS or C++/SC models.

Matchlib “throughput accuracy” is high for:

- Feedforward data streams
- Feedback data streams
- Systems with multiple clocks and multiple data stream rates
- Divergent data streams
- Arbitration of data streams (for any reasonably fair arbitration scheme)
 - e.g. round robin

Design aspects which require extra handling:

- If $II \neq 1$, then need to manually insert waits into pre-HLS model to account for this.
- Currently memory accesses for both internal and external memories are not accounted for in Matchlib timing
 - In other words, Matchlib currently assumes that low level memory accesses themselves are not primary bottlenecks.
 - If they are, then that should be explicitly modeled using Push/Pop or wait() statements.
 - This may be improved in a future release.
- As mentioned earlier, if you use PushNB/PopNB in your DUT or TB, you may introduce timing dependent behavior that can cause pre-HLS vs post-HLS mismatches.
 - But PushNB/PopNB do not introduce timing inaccuracies per se.

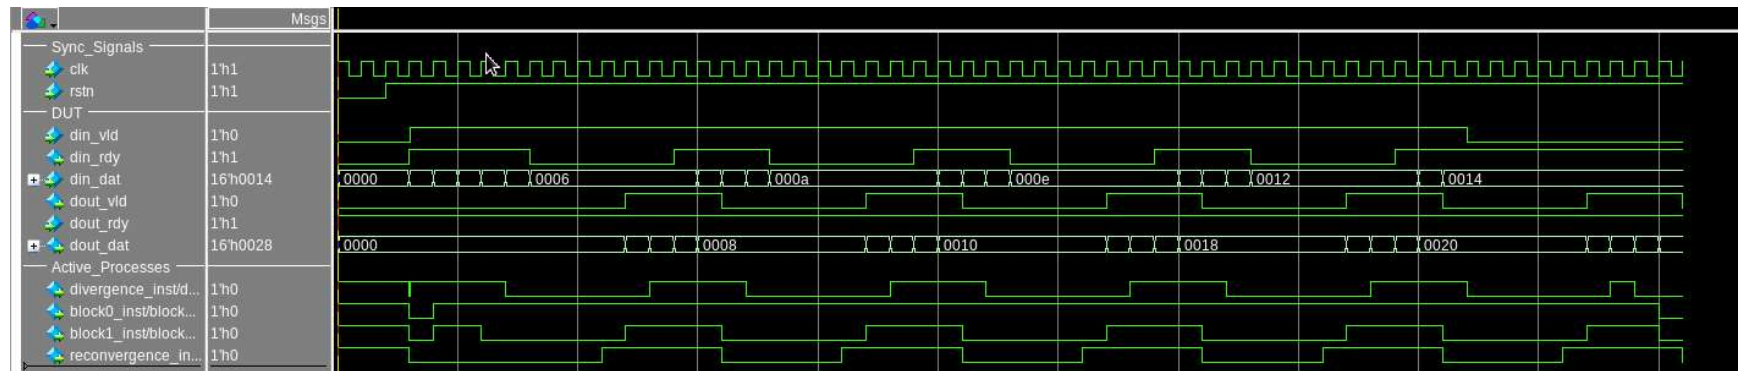

```
void main() {
    out1.Reset();
    in1.Reset();
    wait();
#pragma hls_pipeline_init_interval 4
#pragma pipeline_stall_mode flush
    while (1) {
#ifdef __SYNTHESIS__
        wait();
        wait();
        wait();
#endif
        uint32_t t = in1.Pop();
        out1.Push(t + 0x100);
    }
};
```

How Catapult changes latency and capacity of design...

- When Catapult synthesizes processes/blocks it typically adds latency
 - In other words, post-HLS model will usually have higher latency than pre-HLS model.
- If loop pipelining is used, then it typically also adds capacity.
 - In effect, the HW pipeline is like a HW fifo with the functionality “folded in”.
 - The HW pipeline acts very much like a HW fifo, especially if stall_mode=flush.
- If coupled_io_mode=true is used, then this removes a small amount of capacity and elasticity as compared to pre-HLS model.
- Matchlib supports a “latency and capacity” back-annotation feature to easily back-annotate post-HLS values back into pre-HLS model.
 - This increases accuracy significantly of the pre-HLS sim versus the post-HLS sim.
 - However, you should still not expect exact matches in timing behavior.

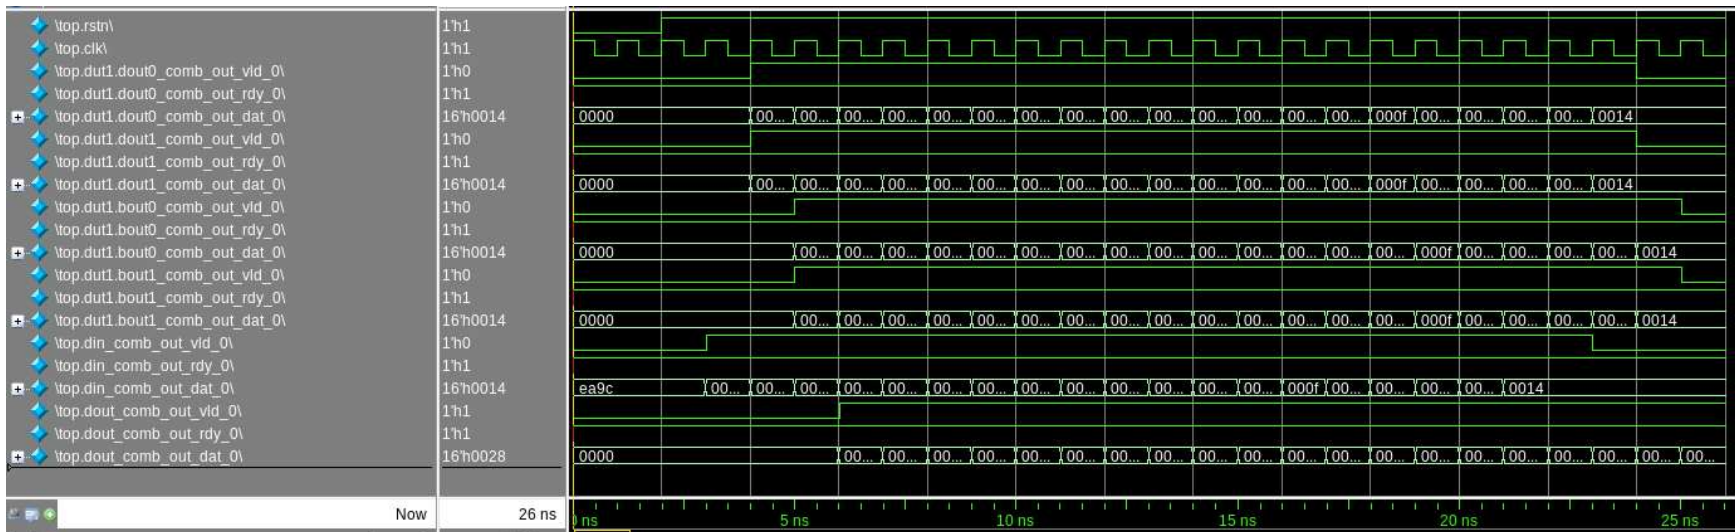

Stuttering Design Example (72* : bagel shop example)

- $II=1$ for all blocks
- Latency of b0 is 1 pre-HLS (by default)
- Latency of b0 is 7 post-HLS
- Post-HLS waveforms below show stuttering due to reconverging data streams with unbalanced latencies and capacities.



Stuttering Design Example Pre-HLS Default Waveforms

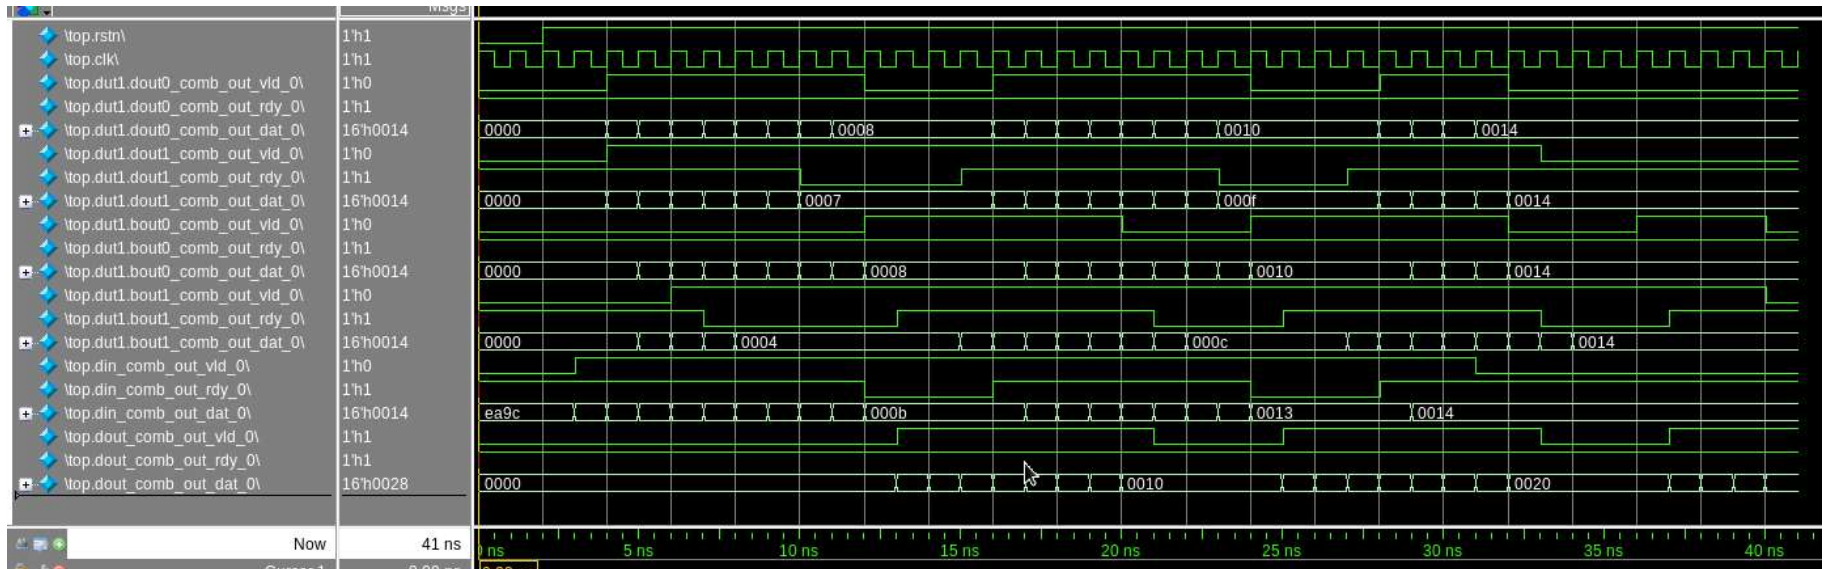
- Latency of b0 is 1 pre-HLS (by default)
- Default pre-HLS waveforms below show no stuttering – new output is seen on every clock.



Stuttering Design Example Pre-HLS with back-annotation

- Now we back-annotate latency and capacity of b0 and b1
- Pre-HLS waveforms below now show stuttering

```
"dut1.bout0_comb_BA": {
  "latency": 7,
  "capacity": 8,
  "src_name": "dut1.block0_inst.dout_vld",
  "dest_name": "dut1.reconvergence_inst.din0_vld"
},
"dut1.bout1_comb_BA": {
  "latency": 1,
  "capacity": 2,
  "src_name": "dut1.block1_inst.dout_vld",
  "dest_name": "dut1.reconvergence_inst.din1_vld"
},
```

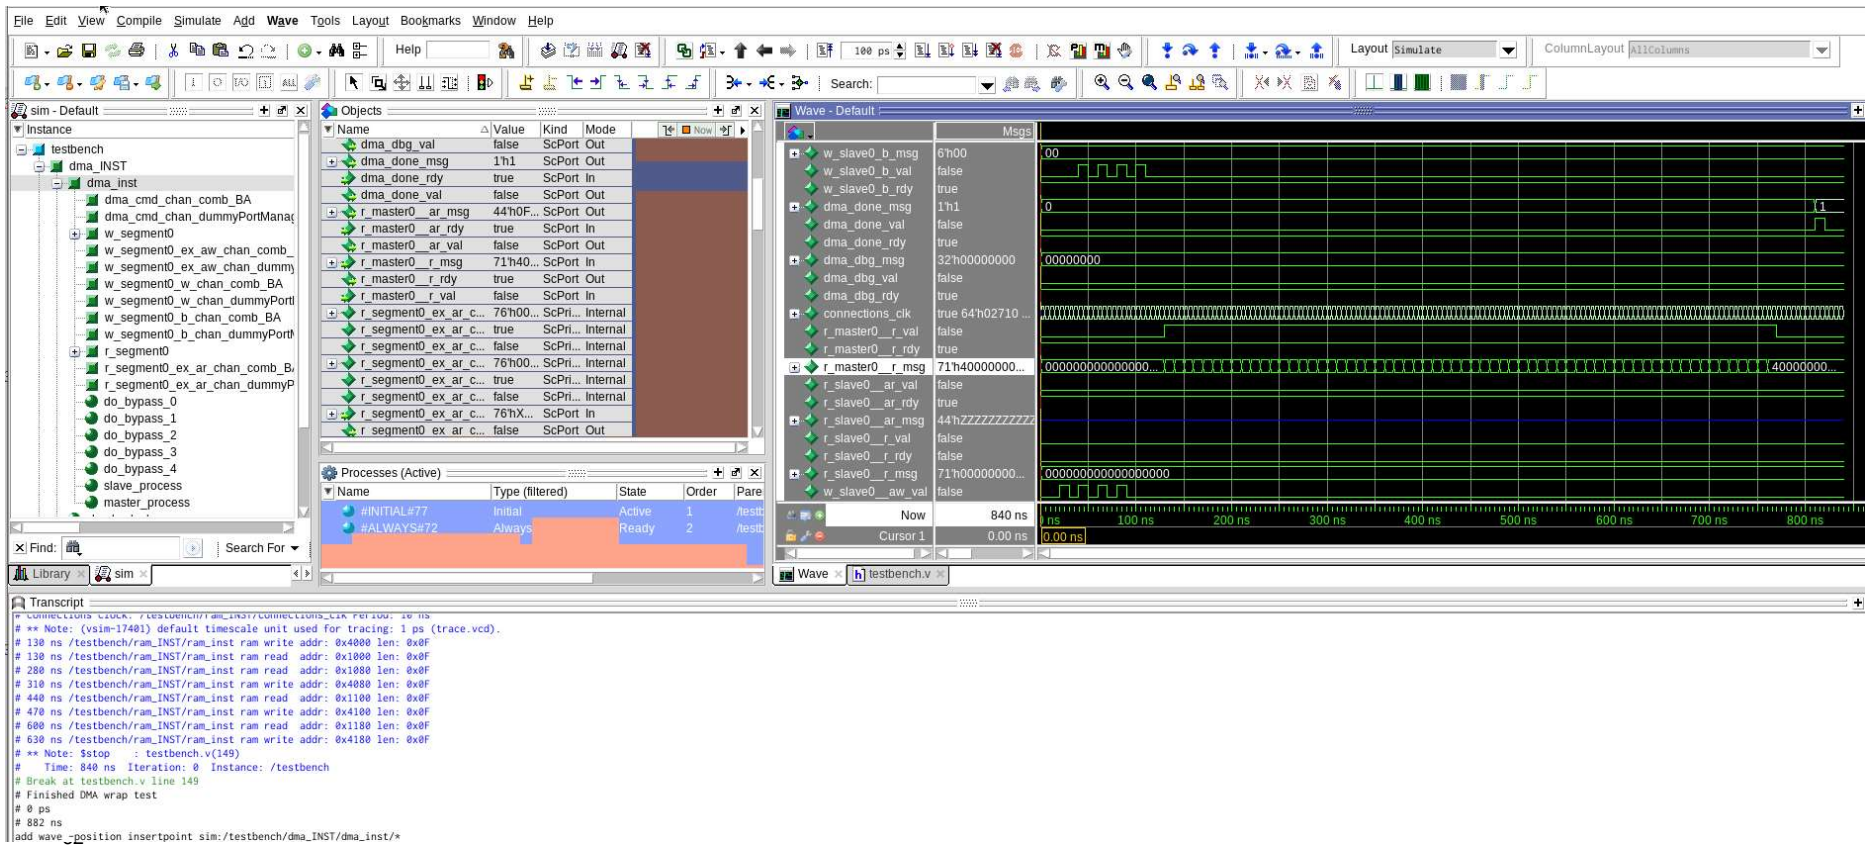


Areas where Matchlib “throughput accuracy” is harder

- Stuttering due to re-convergent data streams (the example just shown)
 - Back annotation can fairly easily show existence of issue and point to the solution
 - Achieving exact match with timing behavior of post-HLS model is harder because:
 - Overall throughput is highly dependent on local latencies and capacities in various blocks
 - Even small things like `coupled_io=true` or `stall_mode=stall` can have a big impact on stuttering in RTL
- Thorough validation that system has no deadlock behaviors
 - Matchlib is useful in finding and resolving many deadlock scenarios in pre-HLS model
 - However, some deadlock scenarios are highly dependent on detailed cycle level behavior and RTL latencies/capacities, and must be verified in the RTL.
- A useful strategy for both of the above scenarios is to “sweep” the latency/capacity values in the pre-HLS back-annotation to stress test the model
 - This will increase the chances that the RTL has no bugs
 - It will also be much easier to debug any issues that are found (as compared to debugging in RTL).

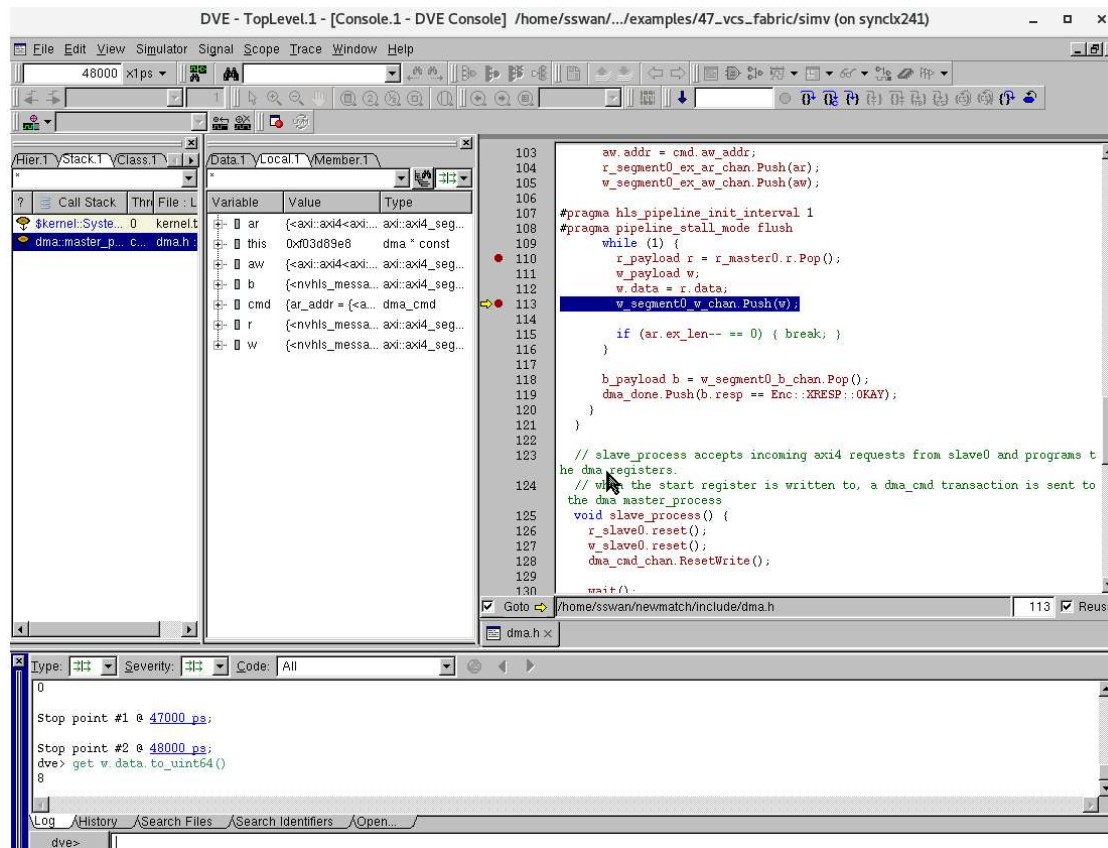
Pre-HLS Debugging in Questa, VCS, Xcelium

- HW – aware debugging in pre-HLS model, can support mixed language (e.g. SV UVM TB + SC DUT)

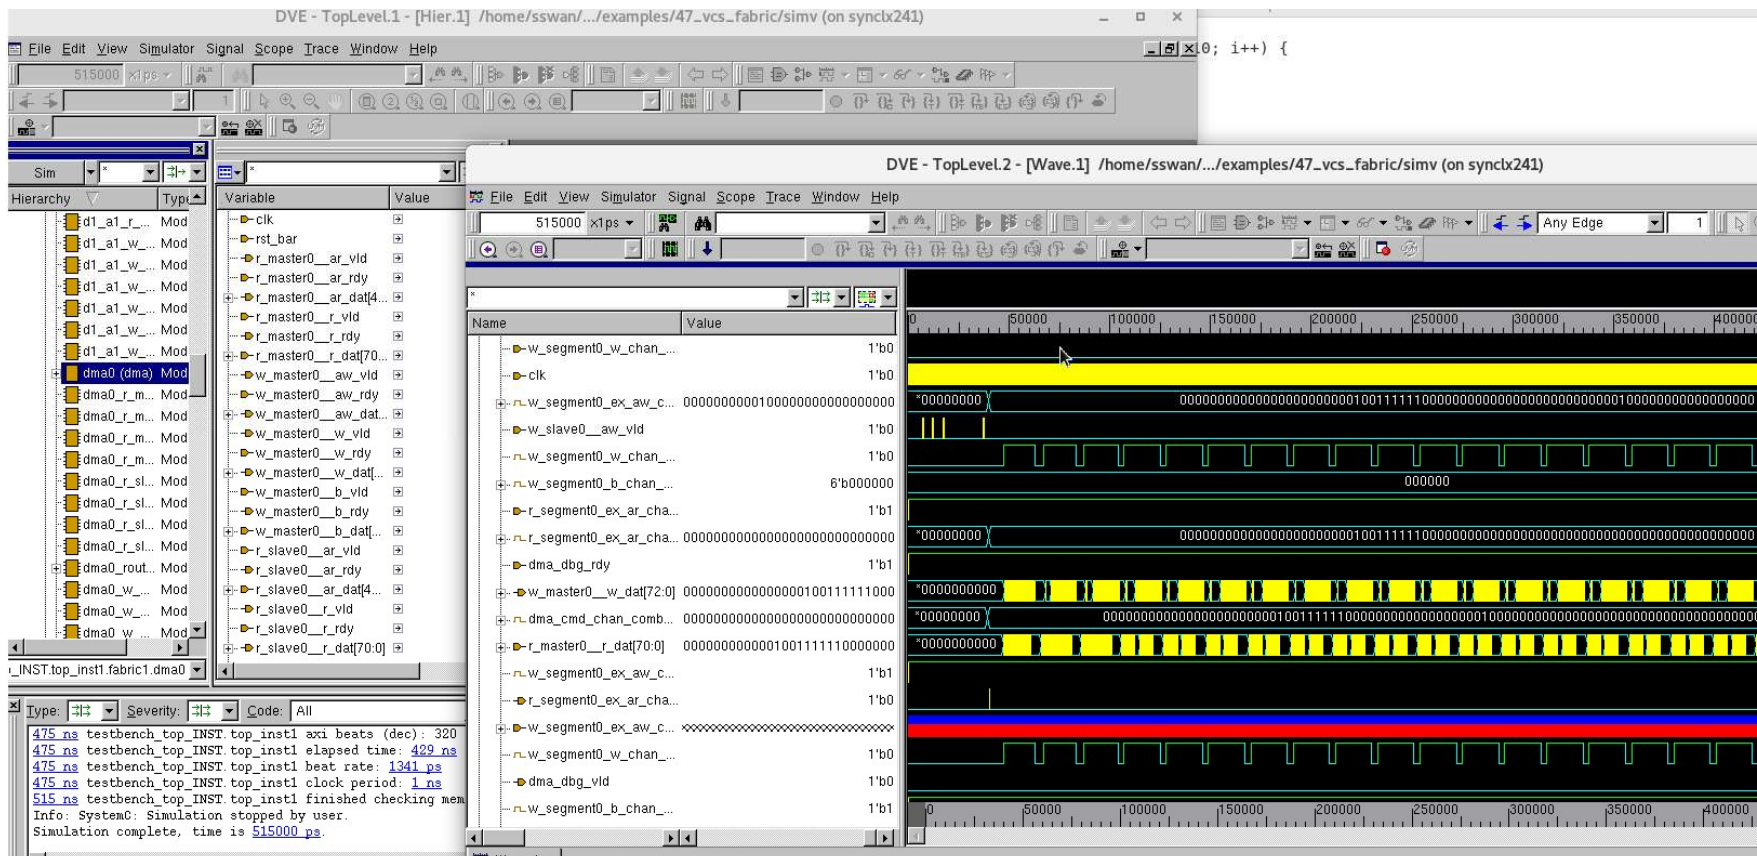


Pre-HLS Debugging in Questa, VCS, Xcelium

- Key advantage: HW aware debugging and throughput accurate modeling with quick edit/debug turnaround..



Pre-HLS Debugging in Questa, VCS, Xcelium (cont)



Supporting Custom Protocols in Catapult SystemC & Matchlib

- First question: Do you really need a new custom protocol?
 - Most modern on-chip bus protocols use message passing (e.g. AXI3/4, NOC protocols, etc).
 - Message passing means "ready/valid/data" signaling.
 - Message passing protocols are easily layered on top of Matchlib Connections
 - Existing AXI4 transactors are built using this approach:
 - See: `$MGC_HOME/shared/examples/matchlib/toolkit/doc/matchlib_customer_training.ppt`
- Second Question: Is it a very simple signal level protocol?
 - Very simple signal level protocols can be easily built by combining Connections::In/Out with sc_signals
 - See `$MGC_HOME/shared/examples/matchlib/toolkit/examples/13_toggle_protocol`

Non-Trivial Custom Protocol Example: ARM APB (52_apb)

- **Step 1:** Study the protocol spec (structural interface and state descriptions)
 - See: <https://developer.arm.com/documentation/ih0024/c>

Signal	Source	Description
PCLK	Clock source	Clock. The rising edge of PCLK times all transfers on the APB.
PRESETn	System bus equivalent	Reset. The APB reset signal is active LOW. This signal is normally connected directly to the system bus reset signal.
PADDR	APB bridge	Address. This is the APB address bus. It can be up to 32 bits wide and is driven by the peripheral bus bridge unit.
PPROT	APB bridge	Protection type. This signal indicates the normal, privileged, or secure protection level of the transaction and whether the transaction is a data access or an instruction access.
PSELx	APB bridge	Select. The APB bridge unit generates this signal to each peripheral bus slave. It indicates that the slave device is selected and that a data transfer is required. There is a PSELx signal for each slave.
PENABLE	APB bridge	Enable. This signal indicates the second and subsequent cycles of an APB transfer.
PWRITE	APB bridge	Direction. This signal indicates an APB write access when HIGH and an APB read access when LOW.
PWDATA	APB bridge	Write data. This bus is driven by the peripheral bus bridge unit during write cycles when PWRITE is HIGH. This bus can be up to 32 bits wide.
PSTRB	APB bridge	Write strobes. This signal indicates which byte lanes to update during a write transfer. There is one write strobe for each eight bits of the write data bus. Therefore, PSTRB[n] corresponds to PWDATA[(8n + 7):(8n)]. Write strobes must not be active during a read transfer.
PREADY	Slave interface	Ready. The slave uses this signal to extend an APB transfer.
PRDATA	Slave interface	Read Data. The selected slave drives this bus during read cycles when PWRITE is LOW. This bus can be up to 32-bits wide.
PSLVERR	Slave interface	This signal indicates a transfer failure. APB peripherals are not required to support the PSLVERR pin. This is true for both existing and new APB peripheral designs. Where a peripheral does not include this pin then the appropriate input to the APB bridge is tied LOW.

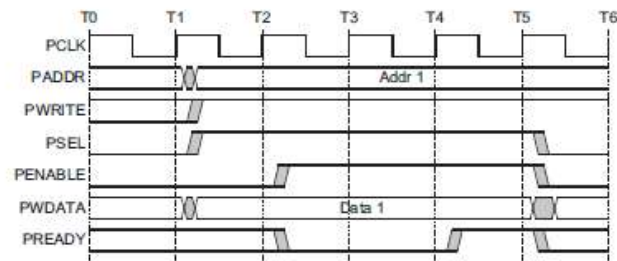


Figure 3-2 Write transfer with wait states

PREADY can take any value when PENABLE is LOW. This ensures that peripherals that have a fixed two cycle access can tie PREADY HIGH.




Figure 3-5 Read transfer with wait states




Figure 4-1 State diagram

Note: FSM perspective is bus, not master or slave. So it is a bit misleading..

Non-Trivial Custom Protocol Example: Step 2

- Abstract the protocol up to the message passing level.
 - Goal: User's HLS model uses message passing interfaces only (Connections::In/Out).
 - All timing details and signal handshakes are hidden inside transactor.
 - We need to define transaction types and the number of message passing interfaces that user's model will see.
 - To get good QOR and be as flexible as possible, the strategy is to keep the message passing interfaces "as close to the HW protocol" as possible while still abstracting away timing behaviors.
- User's HLS model will almost always use blocking message passing calls exclusively.
- Protocol transactor will almost always use non-blocking message passing calls exclusively.
 - It does this while simultaneously managing the signal level protocol and detailed timing.
- Note: Catapult does not support Stratus "protocol regions", and we do not need it to support them.
 - What we do instead is push the "protocol region" into its own thread (in the transactor) and use message passing to communicate with it.

Non-Trivial Custom Protocol Example: Step 2 (cont.)

- For protocols such as AHB which support pipelined behaviors and bursts, it is important to consider how to design the interfaces and transactors to support full throughput.
 - Message passing interfaces make this easier to achieve since they are easily pipelined.
- APB is a simple protocol with no bursts and no simultaneous reads and writes.
- We can simply have a single "req" channel which contains the request, and a single "rsp" channel which contains the response.
 - These two channels are shared between reads and writes because the APB wires for reads and writes are also shared.
 - For a HW protocol which had fully separate wires for reads and writes we would want fully separate wr_req/wr_rsp and rd_req/rd_rsp channels.

Non-Trivial Custom Protocol Example: Step 3

- Create classes to represent APB channels and ports

```
struct apb_sig_chan
{
    apb_sig_chan(const char* name)
        : PSEL(nvhls_concat(name, "_PSEL"))
        , PADDR(nvhls_concat(name, "_PADDR"))
        , PWRITE(nvhls_concat(name, "_PWRITE"))
        , PENABLE(nvhls_concat(name, "_PENABLE"))
        , PWDATA(nvhls_concat(name, "_PWDATA"))
        , PSTRB(nvhls_concat(name, "_PSTRB"))
        , PPROT(nvhls_concat(name, "_PPROT"))
        , PRDATA(nvhls_concat(name, "_PRDATA"))
        , PSLVERR(nvhls_concat(name, "_PSLVERR"))
        , PREADY(nvhls_concat(name, "_PREADY"))
    {}

    SC_SIG(bool,      PSEL);
    SC_SIG(Addr,      PADDR);
    SC_SIG(bool,      PWRITE);
    SC_SIG(bool,      PENABLE);
    SC_SIG(Data,      PWDATA);
    SC_SIG(Wstrb,      PSTRB);
    SC_SIG(Prot_t,     PPROT);
    SC_SIG(Data,      PRDATA);
    SC_SIG(bool,      PSLVERR);
    SC_SIG(bool,      PREADY);
};
```

```
struct apb_master_ports
{
    apb_master_ports(const char* name)
        : PSEL(nvhls_concat(name, "_PSEL"))
        , PADDR(nvhls_concat(name, "_PADDR"))
        , PWRITE(nvhls_concat(name, "_PWRITE"))
        , PENABLE(nvhls_concat(name, "_PENABLE"))
        , PWDATA(nvhls_concat(name, "_PWDATA"))
        , PSTRB(nvhls_concat(name, "_PSTRB"))
        , PPROT(nvhls_concat(name, "_PPROT"))
        , PRDATA(nvhls_concat(name, "_PRDATA"))
        , PSLVERR(nvhls_concat(name, "_PSLVERR"))
        , PREADY(nvhls_concat(name, "_PREADY"))
    {}

    sc_out<bool> PSEL;
    sc_out<Addr> PADDR;
    sc_out<bool> PWRITE;
    sc_out<bool> PENABLE;
    sc_out<Data> PWDATA;
    sc_out<Wstrb> PSTRB;
    sc_out<Prot_t> PPROT;
    sc_in<Data> PRDATA;
    sc_in<bool> PSLVERR;
    sc_in<bool> PREADY;

    template <class C>
    void operator()(C &c) {
        PADDR(c.PADDR);
    }
};
```

```
struct apb_slave_ports
{
    apb_slave_ports(const char* name)
        : PADDR(nvhls_concat(name, "_PADDR"))
        , PWRITE(nvhls_concat(name, "_PWRITE"))
        , PENABLE(nvhls_concat(name, "_PENABLE"))
        , PSEL(nvhls_concat(name, "_PSEL"))
        , PWDATA(nvhls_concat(name, "_PWDATA"))
        , PSTRB(nvhls_concat(name, "_PSTRB"))
        , PPROT(nvhls_concat(name, "_PPROT"))
        , PRDATA(nvhls_concat(name, "_PRDATA"))
        , PSLVERR(nvhls_concat(name, "_PSLVERR"))
        , PREADY(nvhls_concat(name, "_PREADY"))
    {}

    sc_in<Addr> PADDR;
    sc_in<bool> PWRITE;
    sc_in<bool> PENABLE;
    sc_in<bool> PSEL;
    sc_in<Data> PWDATA;
    sc_in<Wstrb> PSTRB;
    sc_in<Prot_t> PPROT;
    sc_out<Data> PRDATA;
    sc_out<bool> PSLVERR;
    sc_out<bool> PREADY;

    template <class C>
    void operator()(C &c) {
        PADDR(c.PADDR);
    }
};
```

Non-Trivial Custom Protocol Example: Step 4

- Define transaction message classes

```
struct apb_req : public nvhls_message {
    bool is_write { false };
    addr_payload addr;
    w_payload w;

    static const unsigned int width = 1 + addr_payload::width;
    template <unsigned int Size> void Marshall(Marshaller<Size> m &r;
        m &is_write;
        m &addr;
        m &w;
    }
    inline friend void sc_trace(sc_trace_file *tf, const
        sc_trace(tf,v.is_write, NAME + ".is_write");
        sc_trace(tf,v.addr, NAME + ".addr");
        sc_trace(tf,v.w, NAME + ".w");
    }
    inline friend std::ostream &operator<<(ostream &os,
        os << rhs.is_write << " ";
        os << rhs.addr << " ";
        os << rhs.w << " ";
        return os;
    }
};
```

```
struct apb_rsp : public nvhls_message {
    r_payload r; // if req was a write, then write resp

    static const unsigned int width = r_payload::width;
    template <unsigned int Size> void Marshall(Marshaller<Size> m &r;
    }
    inline friend void sc_trace(sc_trace_file *tf, const
        sc_trace(tf,v.r, NAME + ".r");
    }
    inline friend std::ostream &operator<<(ostream &os, co
        os << rhs.r << " ";
        return os;
    }
};
```


Non-Trivial Custom Protocol Example: Step 5

- Create the protocol transactor code
 - Strategy: Use PopNB/PushNB to set “got_req” and “got_rsp” flags while simultaneously managing the signal level protocol.

```

230 void main() {
231     // reset all sigs here..
232     wait();
233
234     bool got_req = false;
235     bool got_rsp = false;
236     apb_req req;
237     apb_rsp rsp;
238
239     while (1) {
240         // See ARM APB spec for state description.
241         do {
242             // IDLE state
243             wait();
244             if (!got_req)
245             {
246                 got_req = req_port.PopNB(req);
247             }
248
249             if (got_rsp)
250                 if (rsp_port.PushNB(rsp))
251                 {
252                     got_rsp = 0;
253                 }
254
255         } while (got_rsp || !got_req);
256     }

```

```

190 template <Connections::connections_port_t PortType = AUTO_PORT>
191 class apb_master_xactor : public sc_module {
192 public:
193     sc_in<bool> CCS_INIT_S1(clk);
194     sc_in<bool> CCS_INIT_S1(rst_bar);
195
196     Connections::In<apb_req, PortType> CCS_INIT_S1(req_port);
197     Connections::Out<apb_rsp, PortType> CCS_INIT_S1(rsp_port);
198     sc_out<Addr> CCS_INIT_S1(PADDR);
199     sc_out<bool> CCS_INIT_S1(PWRITE);
    ...

```

```

258 do {
259     wait();
260     if (!got_req)
261     {
262         got_req = req_port.PopNB(req);
263     }
264
265     if (got_rsp)
266         if (rsp_port.PushNB(rsp))
267         {
268             got_rsp = 0;
269         }
270     // SETUP state
271     PSEL = 1;
272     PENABLE = 0;
273     PADDR = req.addr.addr.to_uint64();
274     PWRITE = req.is_write;
275     if (req.is_write) {
276         PWDATA = req.w.data.to_uint64();
277         PSTRB = req.w.wstrb.to_uint64();
278     }
279     else {
280         PWDATA = 0;
281         PSTRB = 0;
282     }
283
284     wait();
285
286     got_req = 0;
287     PSEL = 1;
288     PENABLE = 1;
289
290     do {
291         wait();
292     } while (PREADY == 0);
293     // ACCESS state
294
295     rsp.r.data = PRDATA.read().to_uint64();
296     rsp.r.resp = PSLVERR.read();

```

Non-Trivial Custom Protocol Example: Step 5 (cont.)

- Create the protocol transactor code (apb_slave_transactor here)

```
344 void main() {
345     // reset all sigs here..
346     bool got_req = false;
347     bool got_rsp = false;
348     bool pending_read = false;
349     apb_req req;
350     apb_rsp rsp;
351     wait();
352     while (1) {
353         // See ARM APB spec for state description.
354         do {
355             wait();
356         } while (PSEL == 0);
357         req.is_write = PWRITE.read();
358         req.w.data = PWDATA.read().to_uint64();
359         req.w.wstrb = PSTRB.read().to_uint64();
360         req.addr.addr = PADDR.read().to_uint64();
361         got_req = 1;
362         if (req.is_write)
363             pending_read = false;
364         else
365             pending_read = true;
366         do {
367             if (got_req)
368                 if (req_port.PushNB(req))
369                     {
370                         got_req = 0;
371                     }
372             wait();
373         } while (got_req);
374     }
375 }
```

```
306 template <Connections::connections_port_t PortType = AUTO_PORT>
307 class apb_slave_xactor : public sc_module {
308 public:
309     sc_in<bool> CCS_INIT_S1(clk);
310     sc_in<bool> CCS_INIT_S1(rst_bar);
311     Connections::Out<apb_req, PortType> CCS_INIT_S1(req_port);
312     Connections::In<apb_rsp, PortType> CCS_INIT_S1(rsp_port);
313     sc_in<Addr> CCS_INIT_S1(PADDR);
314     sc_in<bool> CCS_INIT_S1(PWRITE);
315     sc_in<bool> CCS_INIT_S1(PENABLE);
316     sc_in<bool> CCS_INIT_S1(PSEL);
317     sc_in<Data> CCS_INIT_S1(PWDATA);
318     sc_in<Data> CCS_INIT_S1(PSTRB);
319     sc_in<Wstrb> CCS_INIT_S1(PSTRB);
```

```
384     do {
385         wait();
386         if (!got_rsp)
387             {
388                 got_rsp = rsp_port.PopNB(rsp);
389             }
390         while ((PENABLE == 0) || !got_rsp);
391         PREADY = 1;
392         if (pending_read)
393             {
394                 PRDATA = rsp.r.data.to_uint64();
395                 pending_read = false;
396             }
397         PSLVERR = rsp.r.resp; // works for both reads and writes..
398         wait();
399         PREADY = 0;
400         PSLVERR = 0;
401         got_rsp = 0;
402     }
403 }
```

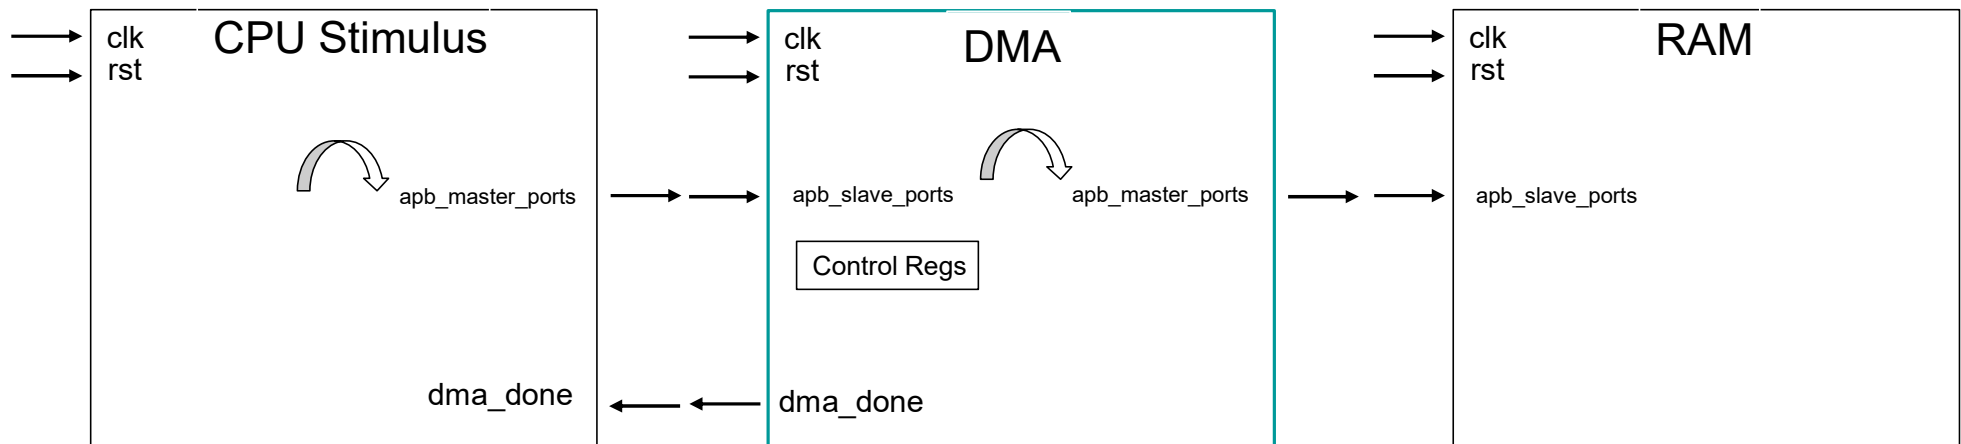
Non-Trivial Custom Protocol Example: Step 6

- Instantiate transactor in user's HLS model

```
50 class dma : public sc_module, public local_apb {
51 public:
52     sc_in<bool> CCS_INIT_S1(clk);
53     sc_in<bool> CCS_INIT_S1(rst_bar);
54
55     apb_master_xactor<> CCS_INIT_S1(master0_xactor);
56     apb_master_ports<> CCS_INIT_S1(master0_ports);
57     apb_req_chan CCS_INIT_S1(master0_req_chan);
58     apb_rsp_chan CCS_INIT_S1(master0_rsp_chan);
59
60     apb_slave_xactor<> CCS_INIT_S1(slave0_xactor);
61     apb_slave_ports<> CCS_INIT_S1(slave0_ports);
62     apb_req_chan CCS_INIT_S1(slave0_req_chan);
63     apb_rsp_chan CCS_INIT_S1(slave0_rsp_chan);
64
65     Connections::Out<bool> CCS_INIT_S1(dma_done);
66     Connections::Out<uint32_t> CCS_INIT_S1(dma_dbg);
67
68     SC_CTOR(dma)
69     {
70         SC_THREAD(slave_process);
71         sensitive << clk.pos();
72         async_reset_signal_is(rst_bar, false);
73
74         SC_THREAD(master_process);
75         sensitive << clk.pos();
76         async_reset_signal_is(rst_bar, false);
77
78         slave0_xactor.clk(clk);
79         slave0_xactor.rst_bar(rst_bar);
80         slave0_xactor.req_port(slave0_req_chan);
81         slave0_xactor.rsp_port(slave0_rsp_chan);
82         slave0_xactor(slave0_ports);
83
84         master0_xactor.clk(clk);
85         master0_xactor.rst_bar(rst_bar);
86         master0_xactor.req_port(master0_req_chan);
87         master0_xactor.rsp_port(master0_rsp_chan);
88         master0_xactor(master0_ports);
89     }
90 }
```

```
99 void master_process() {
100
101     dma_cmd_chan.ResetRead();
102     dma_dbg.Reset();
103     dma_done.Reset();
104
105     master0_req_chan.ResetWrite();
106     master0_rsp_chan.ResetRead();
107
108     wait();
109
110     while(1) {
111         dma_cmd cmd = dma_cmd_chan.Pop();
112         bool status = Enc::XRESP::OKAY;
113         while (1)
114         {
115             apb_req req;
116             apb_rsp rsp;
117
118             req.is_write = false;
119             req.addr.addr = cmd.ar_addr;
120             master0_req_chan.Push(req);
121             rsp = master0_rsp_chan.Pop();
122
123             req.is_write = true;
124             req.addr.addr = cmd.aw_addr;
125             req.w.data = rsp.r.data;
126             master0_req_chan.Push(req);
127             rsp = master0_rsp_chan.Pop();
128
129             if (rsp.r.resp != Enc::XRESP::OKAY)
130                 status = 0;
131
132             if (cmd.len-- == 0)
133                 break;
134
135             cmd.aw_addr += bytesPerBeat;
136             cmd.ar_addr += bytesPerBeat;
137         }
138         dma_done.Push(status);
139     }
```

52_apb Design Example



```
#pragma hls_design top
class dma : public sc_module, public local_apb {
public:
    sc_in<bool> CCS_INIT_S1(clk);
    sc_in<bool> CCS_INIT_S1(rst_bar);

    apb_master_xactor<> CCS_INIT_S1(master0_xactor);
    apb_master_ports<> CCS_INIT_S1(master0_ports);
    apb_req_chan CCS_INIT_S1(master0_req_chan);
    apb_rsp_chan CCS_INIT_S1(master0_rsp_chan);

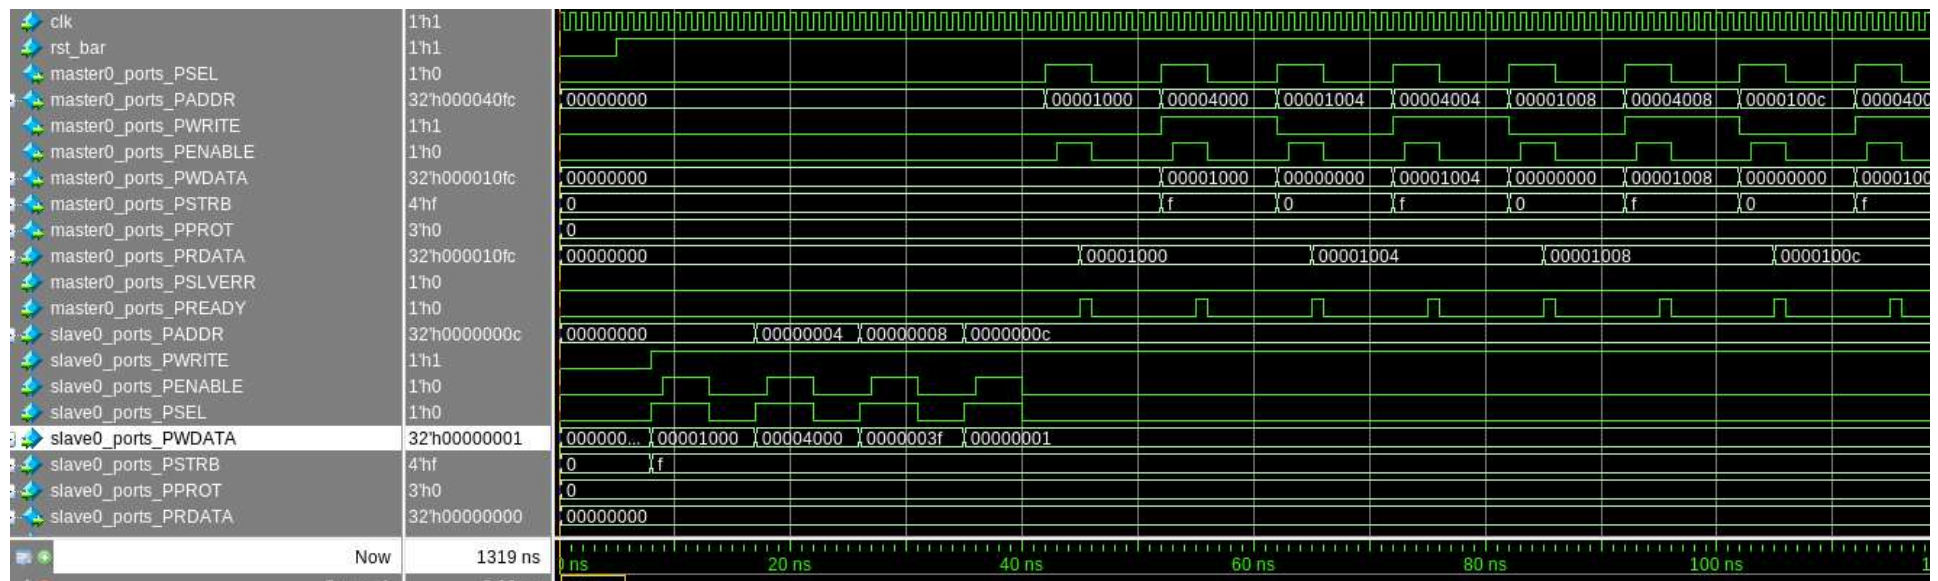
    apb_slave_xactor<> CCS_INIT_S1(slave0_xactor);
    apb_slave_ports<> CCS_INIT_S1(slave0_ports);
    apb_req_chan CCS_INIT_S1(slave0_req_chan);
    apb_rsp_chan CCS_INIT_S1(slave0_rsp_chan);

    Connections::Out<bool> CCS_INIT_S1(dma_done);
    Connections::Out<uint32_t> CCS_INIT_S1(dma_dbg);
};
```

= top level of design

Running 52_apb Example RTL

Note clean preservation of all APB signal names in the Catapult generated Verilog RTL.



Custom Protocol Concluding Thoughts

- If you are more comfortable writing Verilog RTL than SystemC, one approach is to write transactors in Verilog RTL first and then just translate to SC Matchlib afterwards.
 - After all, we are really just writing RTL in the transactors.

Shared Memories in SystemC Designs – Approach #1

Definition: A shared memory is an array in HLS that is preserved thru HLS that is accessed by more than one thread.

Approach #1: “Make the shared memory go away”

- This approach is what many Matchlib users do in practice.
- Use only preserved arrays which are accessed by a **single thread**
- Use architectural directives in Catapult to allocate ports for those rams, etc.
- Use Matchlib components such as Scratchpad, ArbitratedScratchpad, etc., to route requests from other threads to those arrays.
- If throughput is top concern, and a few extra cycles of latency is OK, this is a very robust and flexible approach.
- Pre-HLS sim will be throughput accurate under assumption that mem rd wr access to arrays within threads are not the bottleneck

Shared Memories in SystemC Designs – Approach #2

Approach #2: Use memory model from Catapult Memory Generator

- This is demonstrated in 12_ping_pong_mem
- Each thread accessing memory needs a dedicated RAM port
- There needs to be some synchronization scheme to avoid data races between threads
 - 12_ping_pong_mem uses Connections::SyncChannel
- Catapult memory generator SC models currently need “wait(0.3, SC_NS)” statements deleted from them so they work properly in Matchlib sims
- Catapult memory generator SC models currently do not participate in Matchlib “thruput accurate” simulation mechanism, so sim will not be thruput accurate if thread has > 1 mem accesses per clock

```
43 void thread2() {
44     bool ping_pong = false;
45     out1.Reset();
46     sync1.reset_sync_in();
47     wait();
48
49     #pragma hls_pipeline_init_interval 1
50     #pragma pipeline_stall_mode flush
51     while (1) {
52         sync1.sync_in();
53
54         for (int i=0; i < 8; i++)
55             out1.Push(mem[i + (8 * ping_pong)]);
56
57         ping_pong = !ping_pong;
58     }
59 }
60
61 private:
62     Connections::SyncChannel INIT_S1(sync1); // memory synchronization between threads
63     RAM_IR1W_model<>::mem<ac_int<16>,128> INIT_S1(mem); //Ping-pong shared memory
64 };
```

Memory instance

Memory read operation

Shared Memories in SystemC Designs – Approach #2 (cont.)

Example 12_ping_pong_mem shows shared memory shared by threads in same module.

- In this case, you do not need to explicitly code the mem read and write ports

You can also have memory ports on modules, so that memories can be external and/or shared between multiple modules.

- In this case, you do need to explicitly code the mem read and write ports on module interfaces

```
9 #include "RAM_1R1W.h"
10
11 typedef NVUINTW(32) design_T;
12
13 typedef RAM_1R1W_model<>::mem<design_T,16> mem_t;
14 typedef RAM_1R1W_model<>::rd0_port<design_T,16> mem_rd_t;
15
16
17 #pragma hls_design top
18 class dut : public sc_module {
19 public:
20     sc_in<bool> INIT_S1(clk);
21     sc_in<bool> INIT_S1(rst_bar);
22
23     Connections::Out<NVUINTW(32)> INIT_S1(out1);
24     Connections::In <NVUINTW(32)> INIT_S1(in1);
25     mem_rd_t INIT_S1(mem_rd_port0);
26
27     ac_channel<design_T> chan_in1, chan_out1;
28
29     SC_CTOR(dut)
30     {
31         chan_in1.bind(in1);
32         chan_out1.bind(out1);
33
34         SC_THREAD(main);
35         sensitive << clk.pos();
36         async_reset_signal_is(rst_bar, false);
37     }
38
39 }
```

Memory types

Memory read port

```
13 class Top : public sc_module {
14 public:
15     NVHLS_DESIGN(dut) INIT_S1(dut1);
16
17     sc_clock clk;
18     SC_SIG(bool, rst_bar);
19
20     Connections::Combinational<NVUINTW(32)> INIT_S1(out1);
21     Connections::Combinational<NVUINTW(32)> INIT_S1(in1);
22
23     mem_t INIT_S1(mem0);
24
25     SC_CTOR(Top)
26     : clk("clk", 1, SC_NS, 0.5, 0, SC_NS, true)
27     {
28         Connections::set_sim_clk(&clk);
29         sc_object_tracer<sc_clock> trace_clk(clk);
30
31         dut1.clk(clk);
32         dut1.rst_bar(rst_bar);
33         dut1.out1(out1);
34         dut1.in1(in1);
35         dut1.mem_rd_port0(mem0);
36
37         mem0.CK(clk);
38
39         for (int i=0; i < 16; i++)
40             mem0[i] = i;
41     }
42 }
```

Bind memory read port to memory instance

External Catapult SystemC / Matchlib Resources

Catapult SystemC MatchLib On-Demand Training (ODT)

- <https://eda.learn.sw.siemens.com/training/courses/catapult-high-level-synthesis-library>

2021 Catapult Virtual Seminar (ML Accelerator using Matchlib):

- <https://event.on24.com/wcc/r/3187802/BF0CAE586A768CDB1314C956F64ABA35>

Catapult Matchlib Webinars:

- <https://webinars.sw.siemens.com/how-matchlib-and-systemc-enables>
- <https://webinars.sw.siemens.com/nvidia-design-and-verification-of-a-1>
- <https://webinars.sw.siemens.com/early-axi-soc-performance>

Youtube video from NVidia:

- https://www.youtube.com/watch?v=n8_G-CaSSPU

Accellera SystemC Evolution Day 2020 Matchlib Presentation and Open Source Example Kit:

- <https://forums.accellera.org/files/category/2-systemc/>



Thank you!